

# Crash Course on Deep Learning

Vincent Lepetit


# Class Overview

1. Crash course on Deep Learning;
2. Computer Vision Pipelines and Deep Learning;
3. 3D Localization with Deep Learning;
4. Crash course on Tensor Flow.


Supervised Learning = Regression  
(with (sometimes) some specific properties)

# Regression

$$y = F(x) \quad F?$$



# Linear Regression




Model:  $y = F(x) = ax + b$

Model parameters:  $a$  and  $b$  estimated by optimization, for example:

$$\arg \min_{a,b} \sum_i (ax_i + b - y_i)^2$$

this is called  
"learning" or "training"

# Linear Regression and Prediction



Model:  $y = ax + b$

Now known model parameters:  
 $a$  and  $b$

$$y_{new} = ax_{new} + b$$

this is called "prediction"

# Training Set & Test Set

- Training set:  $\{(x_i, y_i)\}_i$   
known labeled data used to optimize the model's parameters by minimizing the loss function;
- Test set:  $\{(x_i, y_i)\}_i$   
known labeled data used to optimize the model's parameters (and the method in general).  
*The real goal is to have "good" predictions on the test set.*

# Example: Binary Classification



SALMON or SEA BASS?



# Binary Classification


$F(\text{image}) = +1$  for SALMON,  
 $-1$  for SEA BASS




# Deep Networks

- are a possible, general, and now very popular form for  $F(\cdot)$ ;
- are also a general tool for other regression problems.


# A Standard Neural Network



layer's output vector =  
 $\text{sigmoid}(W \cdot \text{input vector} + \text{bias})$




# A Standard Neural Network




final output vector =

$$W_2 \cdot (\text{previous layer's output vector}) + \text{bias}_2$$


# A Standard Neural Network




layer's output vector =  
 $\text{sigmoid}(W \cdot \text{input vector} + \text{bias})$




# The ReLU Operator




layer's output vector =  
 $\text{ReLU}(W \cdot \text{input vector} + \text{bias})$




# Fully Connected Layer



output vector =  $\text{ReLU}(W \cdot \text{input vector} + \text{bias})$




# Fully Connected Layer






# Convolutional Layer



output images =  $\text{ReLU}(\text{filter}_1 * \text{input image},$   
 $\text{filter}_2 * \text{input image}, \dots)$


$$\text{output image}[u, v] = \sum_{i=-m}^{+m} \sum_{j=-n}^n \text{filter}_1[i, j] \text{input image}[u + i, v + j]$$



output images =  $\text{ReLU}(\text{filter}_1 * \text{input image},$   
 $\text{filter}_2 * \text{input image}, \dots)$

# Convolutional Layer


$$\sum_{i=-m}^{+m} \sum_{j=-n}^n \text{filter}_1[i, j] \text{input image}[u + i, v + j]$$



output images =  $\text{ReLU}(\text{filter}_1 * \text{input image},$   
 $\text{filter}_2 * \text{input image}, \dots)$

# Convolutional Layer


$$\sum_{i=-m}^{+m} \sum_{j=-n}^n \text{filter}_1[i, j] \text{input image}[u + i, v + j]$$



output images =  $\text{ReLU}(\text{filter}_1 * \text{input image},$   
 $\text{filter}_2 * \text{input image}, \dots)$

# Convolutional Layer


$$\sum_{i=-m}^{+m} \sum_{j=-n}^n \text{filter}_1[i, j] \text{input image}[u + i, v + j]$$



output images =  $\text{ReLU}(\text{filter}_1 * \text{input image},$   
 $\text{filter}_2 * \text{input image}, \dots)$


# Convolutional Layer

$$\sum_{i=-m}^{+m} \sum_{j=-n}^n \text{filter}_1[i, j] \text{input image}[u + i, v + j]$$



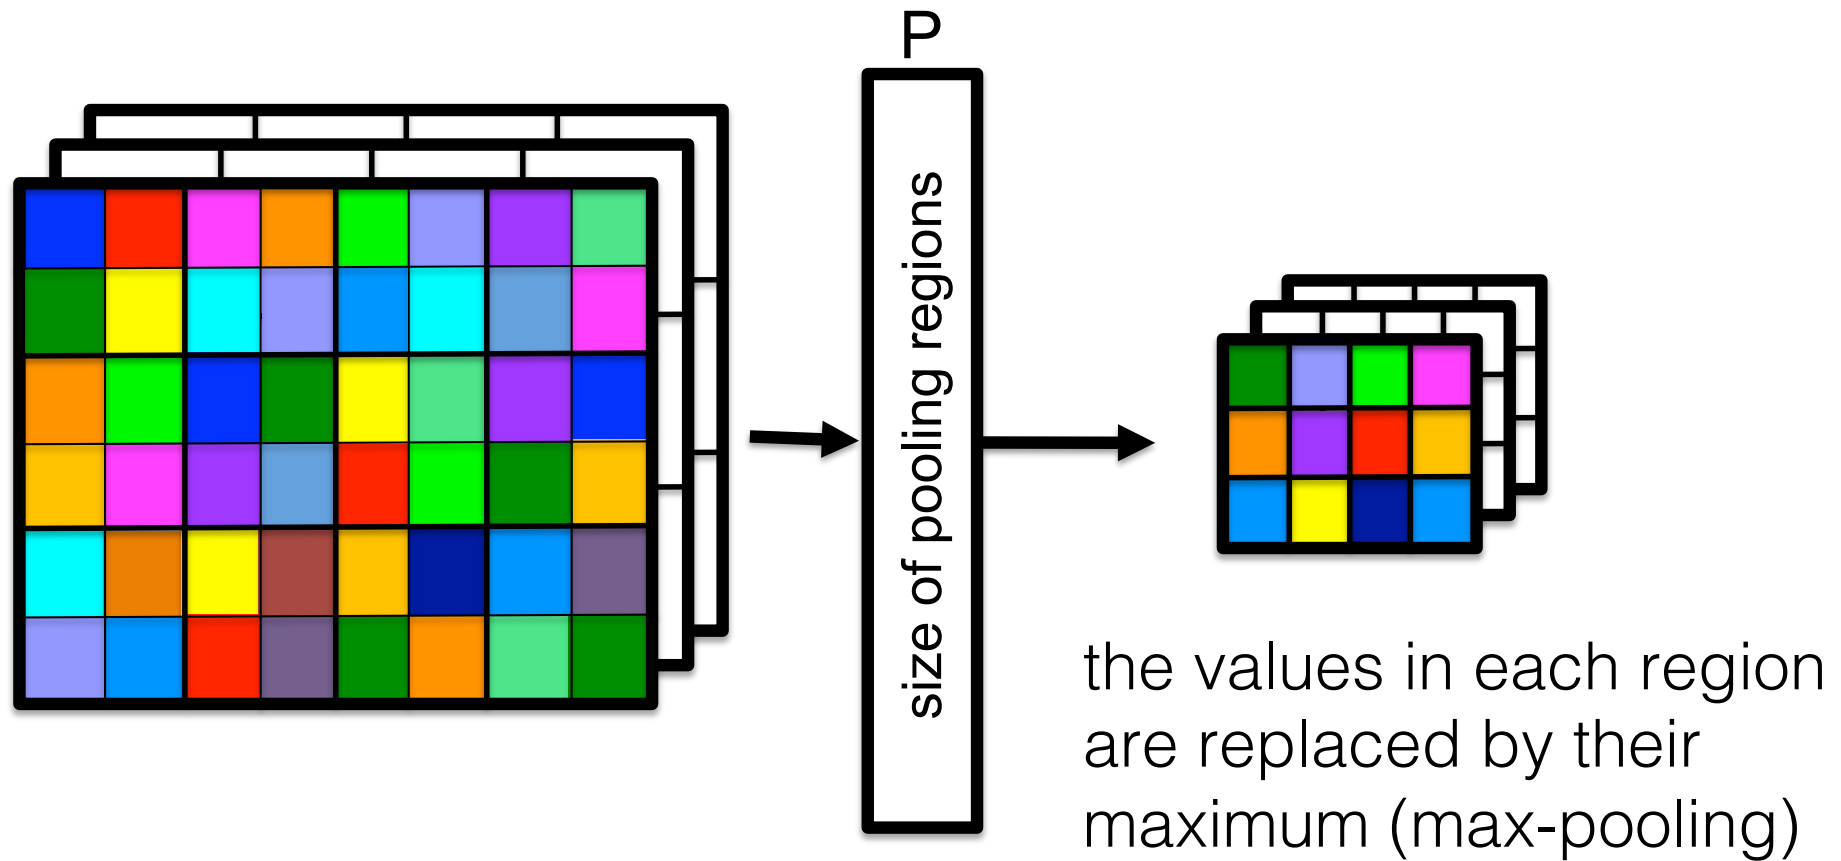
output images =  $\text{ReLU}(\text{filter}_1 * \text{input image},$   
 $\text{filter}_2 * \text{input image}, \dots)$

# Convolutional Layer



output images =  $\text{ReLU}(\text{filter}_1 * \text{input images},$   
 $\text{filter}_2 * \text{input images}, \dots)$

# Pooling Layer






# A Complete Deep Network



Input: An image



$$\mathbf{h}_1 = [\text{ReLU}(f_{1,1} * \mathbf{x}), \dots, \text{ReLU}(f_{1,n} * \mathbf{x})]$$

$$\mathbf{h}_2 = [\text{pool}(\mathbf{h}_{1,1}), \dots, \text{pool}(\mathbf{h}_{1,n})]$$

$$\mathbf{h}_3 = [\text{ReLU}(f_{3,1} * \mathbf{h}_{2,1}), \dots, \text{ReLU}(f_{3,n} * \mathbf{h}_{2,n})]$$

$$\mathbf{h}_4 = [\text{pool}(\mathbf{h}_{3,1}), \dots, \text{pool}(\mathbf{h}_{3,n})]$$

$$\mathbf{h}_5 = \text{ReLU}(\mathbf{W}_5 \mathbf{h}_4 + \mathbf{b}_5)$$


$$\mathbf{h}_6 = \text{ReLU}(\mathbf{W}_6 \mathbf{h}_5 + \mathbf{b}_6)$$

$$\mathbf{h}_7 = \mathbf{W}_7 \mathbf{h}_6 + \mathbf{b}_7$$

$$p(c = \text{car} | I) = \frac{\exp(\mathbf{h}_7[0])}{\exp(\mathbf{h}_7[0]) + \exp(\mathbf{h}_7[1])}$$

$$p(c = \text{no car} | I) = \frac{\exp(\mathbf{h}_7[1])}{\exp(\mathbf{h}_7[0]) + \exp(\mathbf{h}_7[1])}$$

# Optimization



Network parameters found by optimizing an objective function, for example:

$$\min_{\substack{\text{matrices of the fully connected layers} \\ \text{filters of the convolutional layers}}} \sum_{(\text{input}_i, \text{expected output}_i) \text{ in training set}} f(\text{CNN}(\text{input}_i), \text{expected output}_i)$$

Optimization with stochastic gradient descent on mini-batches + dropout + hard example mining + ...

# Two Possible Loss Functions


For multi-class classification, for example using cross-entropy:

$$\min_{\theta} - \sum_i \mathbf{y}_i \log(\text{softmax}_{\theta}(\text{CNN}(\mathbf{x}_i)))$$

For least-squares regression:

$$\min_{\theta} \sum_i \|\mathbf{y}_i - \text{CNN}_{\theta}(\mathbf{x}_i)\|_2^2$$

# My Take on What Deep Networks Do



Consider a 2-layer network of the form:

$$y = \mathbf{w}_2 \sigma(\mathbf{W}_1 \mathbf{x} + \mathbf{b}_1) \quad (1)$$

Introduce the matrix

$$\mathbf{B}(\mathbf{x}) = \text{diag}(\dots, \text{sgn}(\mathbf{W}_1^{(i)} \mathbf{x} + \mathbf{b}_1^{(i)}), \dots)$$


(1) can be rewritten:

$$y = \mathbf{w}_2 \mathbf{B}(\mathbf{x}) (\mathbf{W}_1 \mathbf{x} + \mathbf{b}_1)$$

The function  $\mathbf{x} \mapsto \mathbf{B}(\mathbf{x})$  is piecewise constant.

Thus (1) is piecewise affine (and also continuous).

# 2D Example: Visualizing $F(x)$




Target Function

To be approximated with:

$$y = \mathbf{w}_2 \sigma(\mathbf{W}_1 \mathbf{x} + \mathbf{b}_1)$$

# 2D Example: Visualizing $F(x)$




Target Function

To be approximated with:

$$y = \mathbf{w}_2 \sigma(\mathbf{W}_1 \mathbf{x} + \mathbf{b}_1)$$

$$\mathbf{x} = \begin{bmatrix} u \\ v \end{bmatrix}$$



$$\min_{\mathbf{W}_1, \mathbf{b}_1, \mathbf{w}_2} \sum_i \|\mathbf{w}_2 \sigma(\mathbf{W}_1 \mathbf{x}_i + \mathbf{b}_1) - y_i\|^2$$